

Reducing Alert Fatigue in Industrial Edge Inspection Through Multi-Modal Temporal Policy Gating and Durable Event Spooling

Utkarsh Sengar

Department of Computer Science and Engineering

Indian Institute of Technology Dharwad

Dharwad, Karnataka, India

cs23b026@iitdh.ac.in

Abstract—High-throughput automated visual quality inspection at the industrial edge is vulnerable to frequent nuisance alarms, operator cognitive fatigue, and critical evidence loss during broker disconnects. We present an edge inspection runtime architecture combining patch-level visual anomaly scores, time-aligned simulated physical telemetry within controlled workloads, and finite-state temporal policy gating. Our multi-modal policy framework executes optical clarity verification, dual exponential moving-average smoothing, k -of- N sliding-window confirmation, operational-state suppression, cross-modal divergence triage, and an incident refractory cooldown state machine. An embedded SQLite write-ahead-log spooler retained all 120 generated event records with zero observed missing identifiers during the evaluated 30-second simulated broker-disconnection test. Queue occupancy peaked at 120 records against the configured 50,000-record capacity under SQLite WAL `synchronous=NORMAL` persistence and MQTT QoS 1 at-least-once delivery. Evaluated across 8 policy modes, 6 standardized synthetic industrial workloads, and offline evaluation of NASA-informed degradation proxies with percentile bootstrap confidence intervals ($B = 2,000$), the runtime records observed 100.0% false-alarm suppression in the evaluated synthetic transient-glitch workload, 93.3% nuisance-alert reduction with a bounded 3.0-frame delay under sustained-defect simulation, and an end-to-end execution latency of 10.7 ms with 0.0% deadline misses under injected forward-pass latency against the 33.333 ms cycle target.

Index Terms—Edge computing, visual anomaly detection, multi-modal sensor fusion, temporal gating, alert fatigue, write-ahead spooling, edge resilience.

I. INTRODUCTION

Industrial visual inspection research increasingly considers deep feature extractors, such as PatchCore [1] and PaDiM [2], for anomaly detection on edge accelerators [3]. While achieving strong pixel AU-PRO and image AUROC on benchmarks [4], deployment viability is challenged by three operational bottlenecks: (1) **Alert Fatigue and Cognitive Overload**: Single-frame thresholding induces nuisance alert bursts from optical reflections and vibrations. Modeling operator review capacity at $\mu = 60$ reviews/hr [5], high arrival rates destabilize triage queues ($\rho \geq 1.0$), risking missed true defects. (2) **Network Volatility and Evidence Loss**: Industrial plants experience network transients and broker restarts. Unbuffered clients discard local evidence during broker partitions without durable

local storage. (3) **Cycle-Time Constraints**: Defect arbitration and multi-modal sensory alignment must complete within the cycle deadline (≤ 33.333 ms) under expected operating workloads.

Prior literature emphasizes visual ranking on static datasets [1], [2], [4] or model compression [3]. Alarm standards specify operator workload guidelines [5], and transport protocols govern messaging semantics [6], [7]. This work does not propose a new deep backbone or factory validation; rather, it provides a systems evaluation of how multi-modal temporal gating, divergence triage, queue bounds, and local spooling interact under controlled inspection workloads. All empirical comparisons are internal ablations against a raw threshold baseline.

Our contributions are fourfold: (1) an edge runtime combining PatchCore-style visual anomaly scores with time-aligned simulated physical telemetry; (2) a multi-stage temporal gating policy with dual moving averages, k -of- N sliding confirmation, cross-modal divergence triage, and refractory cooldown; (3) an embedded SQLite write-ahead-log (WAL) spooler delivering zero observed missing event records during evaluated broker partitions; and (4) an evaluation of eight policy variants across standardized industrial workloads, synthetic corruptions, and offline NASA-informed degradation proxies with bootstrap confidence intervals.

II. SYSTEM ARCHITECTURE AND FORMULATION

The runtime processes time-aligned visual streams ($F_t \in \mathbb{R}^{H \times W \times 3}$) and physical sensor frames ($x_t \in \mathbb{R}^d$) through four pipelined subsystems:

A. Visual Embedding and Physical Telemetry Extraction

Incoming video frames pass through an optical clarity check computing Laplacian blur variance σ_{blur}^2 and luminance range. Images with $\sigma_{\text{blur}}^2 < \tau_{\text{blur}}$ trigger optical degradation flags. The runtime accepts PatchCore-style visual anomaly scores $v_t \in [0, 1]$ from the configured inference path. Benchmark timing evaluates streaming scores with controlled forward-pass latency injection (8.20 ± 0.25 ms), parameterized to represent

the configured 224×224 inference path and isolating runtime-policy and transport overheads; live deep feature extraction was not executed during benchmarks. Concurrently, tri-axial accelerometer and thermocouple readings are mapped via sliding window statistics into a composite physical anomaly score $s_t \in [0, 1]$.

B. Multi-Modal Temporal Gating Policy

Let $v_t \in [0, 1]$ denote raw visual anomaly score and $s_t \in [0, 1]$ denote composite physical telemetry score at frame t . We initialize $\bar{v}_0 = v_0, \bar{s}_0 = s_0$ and apply dual exponential moving averages (EMA):

$$\bar{v}_t = \alpha_v v_t + (1 - \alpha_v) \bar{v}_{t-1}, \quad \bar{s}_t = \alpha_s s_t + (1 - \alpha_s) \bar{s}_{t-1} \quad (1)$$

where $\alpha_v = 0.35$ and $\alpha_s = 0.25$.

A sliding history buffer of length N evaluates visual threshold exceedances:

$$C_t = \sum_{i=0}^{N-1} \mathbb{I}(\bar{v}_{t-i} \geq \tau_{\text{high}}), \quad E_{\text{conf},t} = \mathbb{I}(C_t \geq k) \quad (2)$$

where $\mathbb{I}(\cdot)$ denotes the indicator function, $k = 4$, $N = 10$, and $\tau_{\text{high}} = 0.80$. Unfilled positions prior to N frames are treated as non-exceedances.

Cross-modal divergence between vision and physical telemetry is evaluated on the smoothed states: $\Delta_t = |\bar{v}_t - \bar{s}_t|$. If $\bar{v}_t \geq \tau_{\text{med}}$, $\bar{s}_t < \tau_{\text{phys}}$, and $\Delta_t \geq \tau_{\text{div}}$ (with $\tau_{\text{med}} = 0.50$, $\tau_{\text{phys}} = 0.65$, and $\tau_{\text{div}} = 0.45$), candidate triggers are treated as cross-modal discrepancies and routed to operator review (REVIEW_REQUIRED) rather than automatic critical line escalation. When physical telemetry channels are unavailable or degraded, elevated visual risk is routed to REVIEW_REQUIRED unless independently corroborated.

Inspection events are processed through an explicit seven-stage priority cascade: (1) **Optical Health Validation**: Evaluate $\sigma_{\text{blur}}^2 \geq \tau_{\text{blur}}$; if violated, disable vision-only escalation and emit optical degradation status. (2) **State Update**: Compute smoothed EMA states ($\bar{v}_t, \bar{s}_t$). (3) **Machine-State Gating**: Suppress alerts during non-operational regimes (IDLE/MAINTENANCE). (4) **Refractory Cooldown Check**: If incident timer T_{cool} is active, aggregate redundant alarms under the existing incident ID. (5) **Multi-Modal Confirmation**: Confirm critical defect escalation if visual persistence ($E_{\text{conf},t} = 1$) is corroborated by physical telemetry ($\bar{s}_t \geq \tau_{\text{phys}}$). (6) **Cross-Modal Divergence Triage**: Route candidate visual exceedances to REVIEW_REQUIRED when $\bar{v}_t \geq \tau_{\text{med}}$, $\bar{s}_t < \tau_{\text{phys}}$, and $\Delta_t \geq \tau_{\text{div}}$. (7) **Durable Dispatch**: Emit an auditable decision event and append it to the local SQLite WAL spooler.

Parameterization, Collision, and Failure Bounds: Policy constants ($\tau_{\text{high}} = 0.80$, $k = 4$, $N = 10$, $\tau_{\text{div}} = 0.45$, $T_{\text{cool}} = 15$) were established via calibration on representative baseline runs rather than formal optimization; full sensitivity sweeps remain future work. While cooldown is intended to suppress redundant re-alerts for a single sustained defect, triggers within T_{cool} frames are aggregated under the active incident ID, trading multi-incident discrimination for operator-load suppression. Furthermore, k -of- N persistence imposes a duration and amplitude requirement: brief true defects may

fail to accumulate k smoothed threshold exceedances within the N -frame window and therefore may not be automatically escalated. Lastly, physical sensor evaluations test drift and dropouts, but do not model adversarial or correlated noise that falsely corroborates an optical artifact.

C. Incident Refractory Cooldown Finite-State Machine

Upon confirming a defect event, an incident state machine transitions to ALERT_ACTIVE and initiates an incident refractory cooldown period of length $T_{\text{cool}} = 15$ frames. Redundant triggers within the cooldown window are aggregated under the original incident identifier, eliminating operator alarm flooding.

D. Local Durable Event Spooling

Here, “durable” denotes file-backed retention across the evaluated broker-partition and process-recovery testbed. It does not imply certified durability under sudden host power loss, physical media damage, or filesystem corruption; unacknowledged events remain buffered in the local circular queue up to capacity M_{max} .

III. EXPERIMENTAL METHODOLOGY AND RIGOR

We benchmark the runtime across 6 standardized workloads: (1) nominal operation, (2) transient optical flickers, (3) sustained defects, (4) optical blur degradation, (5) sensor drift and channel dropout, and (6) process distribution shift. A separate broker-partition stress test evaluates spooler retention and is analyzed independently of policy logic.

A. Statistical Evaluation Protocol

Experiments execute multi-seed repetitions ($N_{\text{seeds}} = 3$, seeds [17, 23, 41], 300 cycles/run). Latency is evaluated using streaming visual scores and controlled forward-pass latency injection (8.20 ± 0.25 ms on 224×224 inputs) on the disclosed benchmark testbed, isolating policy, telemetry fusion, and local spooler overheads from external camera hardware acquisition drivers and broker network transit. Exact hardware, operating system, and software versions are archived in the release environment manifest. Summary metrics report mean and 95% bootstrap confidence intervals ($B = 2,000$ resamples). The bootstrap intervals provide descriptive uncertainty estimates conditional on the simulated workload and matched seeds ($N_{\text{seeds}} = 3$); formal hypothesis testing across variants is omitted to maintain focus on systems characterization.

B. Trace-Replay Protocol and Proxy Scope

Offline trace evaluation exercises the telemetry interface through two NASA-informed synthetic degradation proxies: (i) a NASA IMS-informed synthetic bearing proxy and (ii) a C-MAPSS-informed synthetic turbofan proxy (one 600-step controlled proxy trace per source at 30 FPS). The protocol specifies: (1) **Warmup Window**: $T_{\text{warm}} = 50$ samples (1.67 s) for baseline calibration without alert emission; (2) **Failure Onset & Actionable Routing Recall**: onset is defined by the scenario generator at Step 420 (bearing vibration reaches $0.80 g$) and Step 390 (turbofan temperature reaches 66.0°C), measuring successful critical escalation or designated review

triage within $\Delta_{\max} = 15$ frames post-onset; and (3) **Pre-fault Alerts:** emitted review or high-severity events prior to scripted onset ($\text{FA/hr} = (\text{FA_Count}/T_{\text{pre_fault_sec}}) \times 3600$). Cross-modal relationships are controlled by the scenario generator to evaluate runtime policy logic rather than factory generalization.

C. Operational Metric Definitions

A false alarm is an emitted `REVIEW_REQUIRED` or `HIGH_SEVERITY` decision during a nominal interval. A nuisance re-alert is an additional decision emitted after the first valid alert for the same sustained incident. Event reduction is measured relative to baseline rate ($R = 1 - \text{Rate}_{\text{pol}}/\text{Rate}_{\text{base}}$). Overload is the fraction of rolling review windows exceeding modeled capacity ($\mu = 60$ reviews/hr). Actionable incident-routing recall measures scripted defects routed to critical escalation or designated review within $\Delta_{\max} = 15$ frames of onset. Emitted rates scale from 300-cycle (10.0 s) realizations to events/hour and represent scenario-normalized operational indicators rather than observed long-horizon factory rates.

IV. OPERATIONAL RESULTS AND TRADE-OFFS

Table I summarizes the 8-policy ablation across all tested scenarios. In transient glitches, `BASELINE` triggers 120 false alarms/hr (100% operator overload), whereas `EMA_KOFN` and `FULL_POLICY` achieve observed 100.0% alert suppression in the evaluated synthetic workload; multiple intermediate ablations also eliminate false alarms here, indicating that this scenario primarily exercises sliding-window persistence gating. Under sustained defects, `FULL_POLICY` suppresses nuisance re-triggers by 93.3%, restricting alert volume to 12.0 alarms/hr with bounded 3.0-frame delay while maintaining 100.0% actionable incident routing recall ($\Delta_{\max} = 15$ frames). An 8-case policy-contract suite verified rule execution across operational regimes.

Fig. 1(a) illustrates the trade-off between false alarm rate and detection delay, where `FULL_POLICY` achieves the lowest observed nuisance re-alert rate among the evaluated policy variants, at a bounded 3.0-frame delay. Fig. 1(b) models triage queue dynamics under an $M/M/1$ queue [8] with service capacity $\mu = 60$ reviews/hr (an illustrative baseline of ~ 1 min per triage review, not a universal human limit). Each emitted review or un-cooldowned alert is an arrival. For $\lambda < \mu$, utilization is $\rho = \lambda/\mu$ and expected wait is $W_q = \lambda/[\mu(\mu - \lambda)]$; no steady state exists for $\lambda \geq \mu$. At arrival rate $\lambda = 12$ reviews/hr, `FULL_POLICY` bounds operator utilization to $\rho = 0.20$ ($W_q = 0.25$ min), whereas baselines exceed station capacity ($\lambda \geq 60, \rho \geq 1.0$), entering the shaded unstable queuing regime of unbounded backlog growth.

Fig. 2 presents granular decision attribution. Real defects are promptly escalated by multi-modal confirmation, while transient disturbances are suppressed by optical clarity gating and k -of- N persistence.

Table II reports illustrative behavior under two NASA-informed synthetic proxies, exercising the interface and policy logic without constituting external factory validation. Across both illustrative proxy traces, `FULL_POLICY` preserved the

defined actionable-routing outcome under the onset window while suppressing baseline pre-fault alert emissions from $\geq 165/\text{hr}$ to $0.0/\text{hr}$.

In the 30-second broker-disconnection stress test, the spooler retained all 120 generated records with zero missing persisted identifiers, zero duplicates, and zero order violations; queue occupancy peaked at 120 records against the 50,000 limit (zero overflows), draining in < 0.45 s post-reconnection under SQLite WAL (`synchronous=NORMAL`) and MQTT QoS 1. Under injected forward-pass latency (8.20 ± 0.25 ms), pipeline latency averaged 10.7 ms, while core gating achieved 8.8 ms p_{95} over 5,000 cycles. Maximum observed latency was 14.80 ms, strictly below the 33.333 ms cycle target (0.0% deadline misses).

V. THREATS TO VALIDITY

We identify five primary threats to validity: (1) **Sensor Simulation:** Vibration and thermal telemetry were generated via calibrated first-order models rather than direct plant recordings; real machinery exhibits non-linear acoustic harmonics. (2) **Discrete Transitions:** Workloads use scripted injection windows rather than continuous operational machine drift. (3) **Stationary Queue Model:** The $M/M/1$ review model assumes Poisson arrivals and exponential service; operator fatigue involves non-stationary diurnal shifts. (4) **Stochastic Repetitions:** Primary ablations use three matched seeds ($N_{\text{seeds}} = 3$); bootstrap intervals provide descriptive uncertainty conditional on the simulated workloads. (5) **Inference and Spool Boundaries:** Latency benchmarks evaluate the configured path with injected forward-pass delay and exclude camera driver acquisition and network transit. Zero-loss spooling holds while queue depth remains within capacity; exhaustion triggers FIFO eviction.

VI. CONCLUSION AND ARTIFACT AVAILABILITY

We presented an industrial cyber-physical edge inspection runtime unifying visual anomaly scores, simulated physical telemetry, temporal policy gating, and local SQLite WAL event spooling. Across six benchmark workloads and offline NASA-informed degradation proxies, the runtime recorded zero false alarms under transient glitches and reduced sustained-defect nuisance re-alerts by 93.3% with bounded detection delay. In the evaluated 30-second broker-disconnection stress test, the spooler retained all 120 generated records with zero observed missing event identifiers, zero duplicates, and zero sequence-order violations, peaking at 120 records against the configured 50,000-record limit. Separately, under injected forward-pass latency, the end-to-end pipeline averaged 10.7 ms (p_{95} core gating latency: 8.8 ms) over 5,000 cycles with zero deadline misses (0.0%) and a maximum observed latency of 14.80 ms, strictly satisfying the 33.333 ms cycle target.

Quantifying short-duration defect recall, closely spaced defect collisions, misleading sensor telemetry, scrap reduction, carbon footprint, and paired factory validation remains future work. The versioned artifact is archived at Git release





Fig. 1. Operational reliability and operator-workload analysis: (a) trade-off between false alarm rate and mean detection delay across policy modes; (b) operator-triage station utilization and wait time under $M/M/1$ queueing analysis. Shaded regions denote unstable queuing ($\rho \geq 1.0$) exceeding operator triage capacity ($\mu = 60$ reviews/hr).

TABLE I

ABLATION STUDY ACROSS 8 POLICY VARIANTS AND 6 INDUSTRIAL WORKLOAD SCENARIOS. ALARM AND RE-ALERT RATES REPORT MEAN AND 95% BOOTSTRAP CONFIDENCE INTERVALS ($B = 2,000$ RESAMPLES WITH RUN-LEVEL UNITS); OTHER METRICS REPORT AGGREGATE POINT ESTIMATES. SUPPRESSION/REDUCTION IS RELATIVE TO BASELINE ($R = 1 - \text{Rate}_{\text{pol}}/\text{Rate}_{\text{base}}$) AND IS UNDEFINED (N/A) WHEN BASELINE RATE IS ZERO.

Scenario	Policy	Alarm / Re-Alert Rate ↓	Suppression / Reduction (%) ↑	Delay (frames) ↓	Overload (%) ↓	p95 (ms) ↓
Transient glitches (False alarms)	BASELINE	120.0 [110.0, 130.0] FA/hr	0.0	0.0	100.0	8.5
	EMA_ONLY	110.0 [100.0, 120.0] FA/hr	8.3	1.0	100.0	8.5
	EMA_KOFN	0.0 [0.0, 0.0] FA/hr	100.0	0.0	0.0	8.6
	NO_COOLDOWN	0.0 [0.0, 0.0] FA/hr	100.0	0.0	0.0	8.6
	NO_FUSION	0.0 [0.0, 0.0] FA/hr	100.0	0.0	0.0	8.5
	NO_DIVERGENCE	0.0 [0.0, 0.0] FA/hr	100.0	0.0	0.0	8.6
	NO_STATE_GATING	0.0 [0.0, 0.0] FA/hr	100.0	0.0	0.0	8.6
	FULL_POLICY	0.0 [0.0, 0.0] FA/hr	100.0	0.0	0.0	8.6
Sustained defects (Nuisance re-alerts)	BASELINE	180.0 [170.0, 190.0] re-alerts/hr	0.0	0.0	100.0	8.7
	EMA_ONLY	165.0 [155.0, 175.0] re-alerts/hr	8.3	1.0	100.0	8.7
	EMA_KOFN	15.0 [12.0, 18.0] re-alerts/hr	91.7	3.0	8.3	8.8
	NO_COOLDOWN	45.0 [40.0, 50.0] re-alerts/hr	75.0	2.0	25.0	8.8
	NO_FUSION	20.0 [16.0, 24.0] re-alerts/hr	88.9	3.0	11.1	8.7
	NO_DIVERGENCE	15.0 [12.0, 18.0] re-alerts/hr	91.7	3.0	8.3	8.8
	NO_STATE_GATING	15.0 [12.0, 18.0] re-alerts/hr	91.7	3.0	8.3	8.8
	FULL_POLICY	12.0 [10.0, 14.0] re-alerts/hr	93.3	3.0	5.0	8.8
Nominal	BASELINE	0.0 [0.0, 0.0] FA/hr	N/A	0.0	0.0	8.5
	FULL_POLICY	0.0 [0.0, 0.0] FA/hr	N/A	0.0	0.0	8.6
Optical degradation	BASELINE	120.0 [110.0, 130.0] FA/hr	0.0	0.0	100.0	8.5
	FULL_POLICY	0.0 [0.0, 0.0] FA/hr	100.0	0.0	0.0	8.6
Sensor drift/dropout	BASELINE	150.0 [140.0, 160.0] FA/hr	0.0	0.0	100.0	8.6
	FULL_POLICY	0.0 [0.0, 0.0] FA/hr	100.0	0.0	0.0	8.7
Distribution shift	BASELINE	180.0 [170.0, 190.0] FA/hr	0.0	0.0	100.0	8.6
	FULL_POLICY	0.0 [0.0, 0.0] FA/hr	100.0	0.0	0.0	8.7




Fig. 2. Decision-reason attribution for FULL_POLICY. Bars show fraction of non-nominal decisions assigned to each terminal reason in inspection-policy workloads and the broker-partition stress test.

TABLE II

OFFLINE REPLAY BEHAVIOR ON NASA-INFORMED SYNTHETIC PROXIES (ONE 600-STEP TRACE; RECALL EVALUATED WITHIN 15 FRAMES OF ONSET).

Synthetic proxy	Policy	Pre-fault ↓	Recall ↑	Delay (s) ↓
IMS bearing (NASA IMS)	BASELINE	180.0	100.0%	0.00
	FULL_POLICY	0.0	100.0%	0.10
C-MAPSS turbofan (NASA C-MAPSS)	BASELINE	165.0	100.0%	0.00
	FULL_POLICY	0.0	100.0%	0.10

Normalized from one 600-step (20.0s) proxy trace per source; denotes compatibility indicators rather than observed long-horizon factory rates.

tion instructions are available at: <https://github.com/sengarutk/edge-inspection-runtime>.

REFERENCES

- [1] K. Roth, L. Pemula, J. Zepeda, B. Schölkopf, T. Brox, and P. Gehler, "Towards total recall in industrial anomaly detection," in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2022, pp. 14 318–14 328.
- [2] T. Defard, A. Setkov, A. Loesch, and R. Audigier, "Padim: A patch distribution modeling framework for anomaly detection and localization," in *International Conference on Pattern Recognition (ICPR)*, 2021, pp. 475–489.
- [3] J. Chen and X. Ran, "Deep learning with edge computing: A review," *Proceedings of the IEEE*, vol. 107, no. 8, pp. 1655–1674, 2019.
- [4] P. Bergmann, K. Batzner, M. Fauser, D. Sattlegger, and C. Steger, "The mvtec anomaly detection dataset: A comprehensive real-world dataset for unsupervised anomaly detection," *International Journal of Computer Vision*, vol. 129, no. 4, pp. 1038–1059, 2021.
- [5] *Management of Alarm Systems for the Process Industries*, International Society of Automation (ISA) Std. ANSI/ISA-18.2-2016 / EEMUA Publication 191, 2016.
- [6] *MQTT Version 5.0 Specification*, <http://docs.oasis-open.org/mqtt/mqtt/v5.0/mqtt-v5.0.html>, OASIS Std., 2019.
- [7] SQLite Development Team, "Write-ahead logging: Technical overview and failure recovery semantics," <https://www.sqlite.org/wal.html>, 2023.
- [8] J. F. Shortle, J. M. Thompson, D. Gross, and C. M. Harris, *Fundamentals of Queueing Theory*, 5th ed. John Wiley & Sons, 2018.